

Assignment 4

Due at 11:59pm on November 5.

Yao Sun

This is an individual assignment. Turn in this assignment as an HTML or PDF file to ELMS. Make sure to include the R Markdown or Quarto file that was used to generate it. Include the GitHub link for the repository containing these files.

- The name on the README.md is homework4. The link to the GitHub repository is <https://github.com/YSun-umich/homework4.git>

In this notebook we will use Google BigQuery, “Google’s fully managed, petabyte scale, low cost analytics data warehouse”. Some instruction on how to connect to Google BigQuery can be found here: <https://db.rstudio.com/databases/big-query/>.

You will need to set up a Google account with a project to be able to use this service. We will be using a public dataset that comes with 1 TB/mo of free processing on Google BigQuery. As long as you do not repeat the work in this notebook constantly, you should be fine with just the free tier.

Go to <https://console.cloud.google.com> and make sure you are logged in a non-university Google account. **This may not work on a university G Suite account because of restrictions on those accounts.** Create a new project by navigating to the dropdown menu at the top (it might say “Select a project”) and selecting “New Project” in the window that pops up. Name it something useful.

After you have initialized a project, paste your project ID into the following chunk.

```
project <- "yaosun-survm727-hw4"
```

We will connect to a public database, the Chicago crime database, which has data on crime in Chicago.

```
con <- dbConnect(  
  bigrquery::bigquery(),  
  project = "bigquery-public-data",  
  dataset = "chicago_crime",  
  billing = project  
)  
  
con
```

```
<BigQueryConnection>  
  Dataset: bigquery-public-data.chicago_crime  
  Billing: yaosun-survm727-hw4
```

We can look at the available tables in this database using `dbListTables`.

Note: When you run this code, you will be sent to a browser and have to give Google permissions to Tidyverse API Packages. **Make sure you select all to give access or else your code will not run.**

```
dbListTables(con)
```

```
i Suitable tokens found in the cache, associated with these emails:
```

```
* 'sunyao2021@gmail.com'
```

```
* 'sunyao@umich.edu'
```

```
Defaulting to the first email.
```

```
! Using an auto-discovered, cached token.
```

```
To suppress this message, modify your code or options to clearly consent to  
the use of a cached token.
```

```
See gargle's "Non-interactive auth" vignette for more details:
```

```
<https://gargle.r-lib.org/articles/non-interactive-auth.html>
```

```
i The bigrquery package is using a cached token for 'sunyao2021@gmail.com'.
```

```
[1] "crime"
```

Information on the ‘crime’ table can be found here:

<https://cloud.google.com/bigquery/public-data/chicago-crime-data>

Write a first query that counts the number of rows of the ‘crime’ table in the year 2016. Use code chunks with {sql connection = con} in order to write SQL code within the document.

```
SELECT count(primary_type), count(*)
FROM crime
WHERE year = 2016
LIMIT 10;
```

Table 1: 1 records

f0__	f1__
269922	269922

Next, count the number of arrests grouped by `primary_type` in 2016. Note that is a somewhat similar task as above, with some adjustments on which rows should be considered. Sort the results, i.e. list the number of arrests in a descending order.

```
SELECT primary_type, count(*) AS arrest_count
FROM crime
WHERE EXTRACT(YEAR FROM date) = 2016 AND arrest = TRUE
GROUP BY primary_type
ORDER BY arrest_count DESC;
```

Table 2: Displaying records 1 - 10

primary_type	arrest_count
NARCOTICS	13327
BATTERY	10333
THEFT	6522
CRIMINAL TRESPASS	3724
ASSAULT	3492
OTHER OFFENSE	3415
WEAPONS VIOLATION	2511
CRIMINAL DAMAGE	1669

primary_type	arrest_count
PUBLIC PEACE VIOLATION	1116
MOTOR VEHICLE THEFT	1098

We can also use the `date` for grouping. Count the number of arrests grouped by hour of the day in 2016. You can extract the latter information from `date` via `EXTRACT(HOUR FROM date)`. Which time of the day is associated with the most arrests?

```
SELECT EXTRACT(HOUR FROM date) AS hour_of_day, COUNT(*) AS arrest_count
FROM crime
WHERE EXTRACT(YEAR FROM date) = 2016 AND arrest = TRUE
GROUP BY hour_of_day
ORDER BY arrest_count DESC;
```

Table 3: Displaying records 1 - 10

hour_of_day	arrest_count
19	3843
18	3481
20	3302
21	2961
16	2933
22	2896
11	2895
17	2820
12	2787
14	2774

Focus only on `HOMICIDE` and count the number of arrests for this incident type, grouped by year. List the results in descending order.

```
SELECT EXTRACT(YEAR FROM date) AS Year, COUNT(*) as Homicide_Arrests_Count
FROM crime
WHERE primary_type = 'HOMICIDE' AND arrest = true
GROUP BY Year
ORDER BY Homicide_Arrests_Count DESC;
```

Table 4: Displaying records 1 - 10

Year	Homicide_Arrests_Count
2001	430
2002	427
2003	382
2020	349
2022	306
2004	294
2021	292
2016	289
2008	287
2005	284

Find out which districts have the highest numbers of arrests in 2015 and 2016. That is, count the number of arrests in 2015 and 2016, grouped by year and district. List the results in descending order.

```
SELECT EXTRACT(YEAR FROM date) AS year, district, COUNT(*) AS arrest_count
FROM crime
WHERE EXTRACT(YEAR FROM date) IN (2015, 2016) AND arrest = TRUE
GROUP BY year, district
ORDER BY year DESC, arrest_count DESC;
```

Table 5: Displaying records 1 - 10

year	district	arrest_count
2016	11	6575
2016	7	3655
2016	6	3449
2016	15	3073
2016	10	2951
2016	8	2951
2016	25	2950
2016	4	2838
2016	5	2702
2016	9	2592

Lets switch to writing queries from within R via the DBI package. Create a query object that counts the number of arrests grouped by **primary_type** of district 11 in year 2016. The results should be displayed in descending order.

Execute the query.

```
query <- "SELECT primary_type, COUNT(*) AS arrest_count
FROM crime
WHERE EXTRACT(YEAR FROM date) = 2016 AND arrest = TRUE AND district = 11
GROUP BY primary_type
ORDER BY arrest_count DESC"

dbGetQuery(con, query)
```

```
# A tibble: 27 x 2
  primary_type      arrest_count
  <chr>            <int>
1 NARCOTICS        3634
2 BATTERY          635
3 PROSTITUTION     511
4 WEAPONS VIOLATION 303
5 OTHER OFFENSE    255
6 ASSAULT          206
7 CRIMINAL TRESPASS 205
8 PUBLIC PEACE VIOLATION 135
9 INTERFERENCE WITH PUBLIC OFFICER 119
10 CRIMINAL DAMAGE  106
# i 17 more rows
```

Try to write the very same query, now using the `dbplyr` package. For this, you need to first map the `crime` table to a tibble object in R.

```
crime_query <- "SELECT * FROM crime"

crime_tibble <- tbl(con, sql(crime_query))
```

Again, count the number of arrests grouped by `primary_type` of district 11 in year 2016, now using `dplyr` syntax.

```
crime_data <- crime_tibble %>% filter(year == 2016 & arrest == TRUE & district == 11) %>% gr
crime_data
```

```
# Source:      SQL [?? x 2]
# Database:    BigQueryConnection
```

```
# Ordered by: desc(arrest_count)
  primary_type      arrest_count
  <chr>             <int>
1 NARCOTICS        3634
2 BATTERY          635
3 PROSTITUTION     511
4 WEAPONS VIOLATION 303
5 OTHER OFFENSE    255
6 ASSAULT          206
7 CRIMINAL TRESPASS 205
8 PUBLIC PEACE VIOLATION 135
9 INTERFERENCE WITH PUBLIC OFFICER 119
10 CRIMINAL DAMAGE 106
# i more rows
```

Count the number of arrests grouped by `primary_type` and `year`, still only for district 11. Arrange the result by `year`.

```
crime_data_2 <- crime_tibble %>% filter(district == 11, arrest == TRUE) %>% group_by(primary_type, year)
crime_data_2
```

``summarise()`` has grouped output by "primary_type". You can override using the ``.groups`` argument.

```
# Source:      SQL [?? x 3]
# Database:    BigQueryConnection
# Groups:      primary_type
# Ordered by:  desc(arrest_count)
  primary_type  year arrest_count
  <chr>         <int>      <int>
1 NARCOTICS     2005        9718
2 NARCOTICS     2003        9562
3 NARCOTICS     2002        9232
4 NARCOTICS     2004        9083
5 NARCOTICS     2006        8185
6 NARCOTICS     2001        7979
7 NARCOTICS     2007        7395
8 NARCOTICS     2013        7234
9 NARCOTICS     2014        6801
10 NARCOTICS    2009        5942
# i more rows
```

Assign the results of the query above to a local R object.

```
crime_data <- crime_data %>% collect()
crime_data_2 <- crime_data_2 %>% collect()
```

`summarise()` has grouped output by "primary_type". You can override using the `groups` argument.

Confirm that you pulled the data to the local environment by displaying the first ten rows of the saved data set.

```
head(crime_data, 10)
```

```
# A tibble: 10 x 2
  primary_type      arrest_count
  <chr>            <int>
1 NARCOTICS        3634
2 BATTERY          635
3 PROSTITUTION     511
4 WEAPONS VIOLATION 303
5 OTHER OFFENSE    255
6 ASSAULT          206
7 CRIMINAL TRESPASS 205
8 PUBLIC PEACE VIOLATION 135
9 INTERFERENCE WITH PUBLIC OFFICER 119
10 CRIMINAL DAMAGE  106
```

```
head(crime_data_2, 10)
```

```
# A tibble: 10 x 3
# Groups:   primary_type [1]
  primary_type year arrest_count
  <chr>      <int>      <int>
1 NARCOTICS   2005      9718
2 NARCOTICS   2003      9562
3 NARCOTICS   2002      9232
4 NARCOTICS   2004      9083
5 NARCOTICS   2006      8185
6 NARCOTICS   2001      7979
7 NARCOTICS   2007      7395
```


8	NARCOTICS	2013	7234
9	NARCOTICS	2014	6801
10	NARCOTICS	2009	5942

Close the connection.

```
dbDisconnect(con)
```